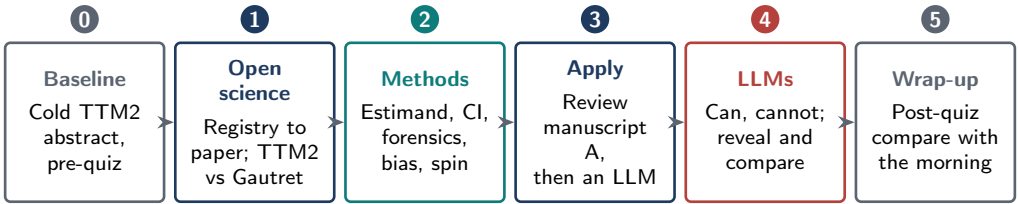
Critical Appraisal of Scientific Papers

Reading a clinical paper the way a good peer reviewer does, with and without AI

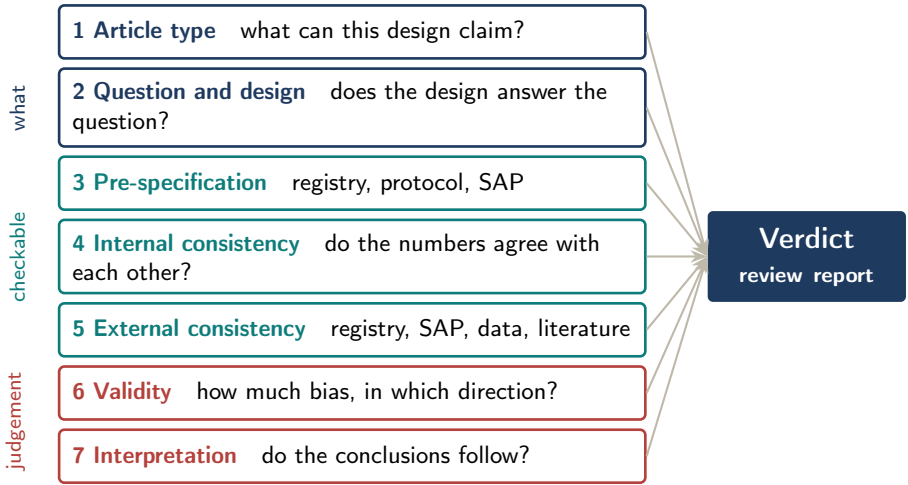
Imad El Badisy

One-day workshop

The day: learn on two real papers, apply on one

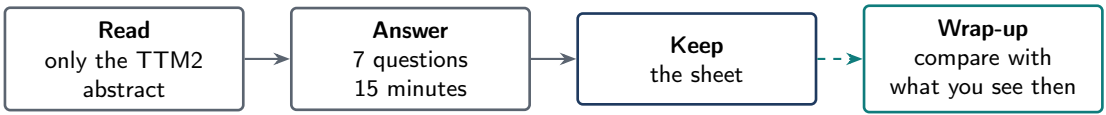


One tool all day: seven layers, then a verdict



Checklist: handouts/checklist-7-layers.md

Cold appraisal, before any teaching



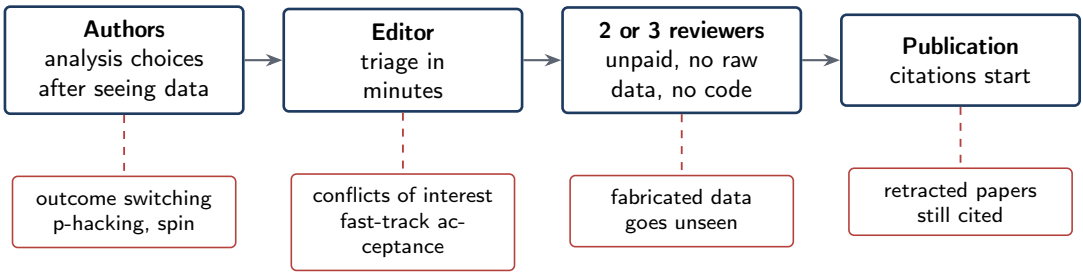
Exercise 00 individual, no AI, no internet

Is TTM2 a "negative" trial? Does it show that hypothermia does not work? What would you check before trusting it?

1 Open science and pre-specification

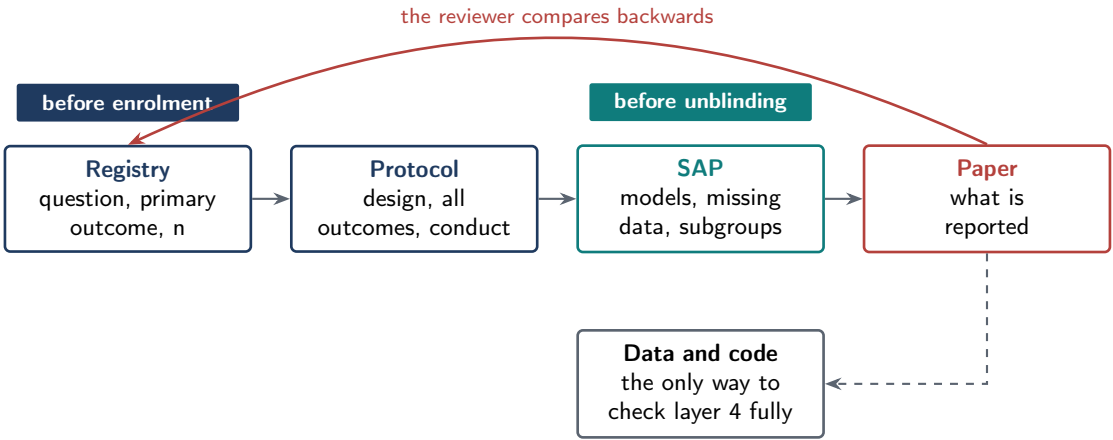
What was promised, what was published, and who can check. Layers 1, 2, 3 and 5.

Where publishing breaks



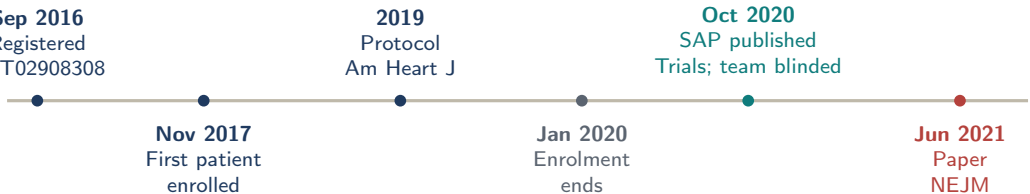
Cases used today: Gautret 2020 (fast-track, retracted 2024); Elgazzar 2020 (withdrawn, fabricated records).

Four documents, each locks something before the data are seen



SPIRIT 2025 (protocols), CONSORT 2025 (reports), ICH E9(R1) (estimands and SAP).

TTM2: a paper trail you can follow

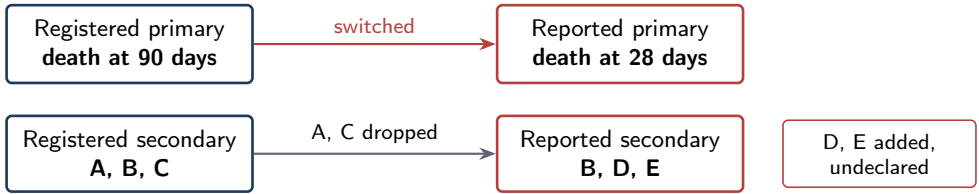


Primary outcome
death at 6 months:
same in all four documents

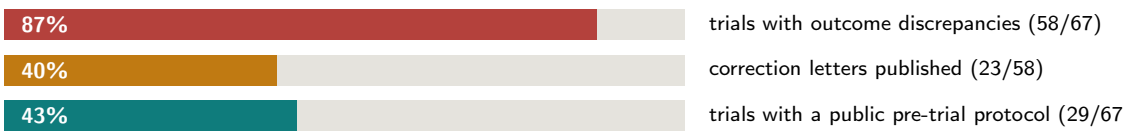
Sample size 1200 to 1900
for TAME co-enrolment,
before any results

One detail
SAP relative reduction 13.6%,
paper says "15%"

Outcome switching: the registered outcome quietly becomes another

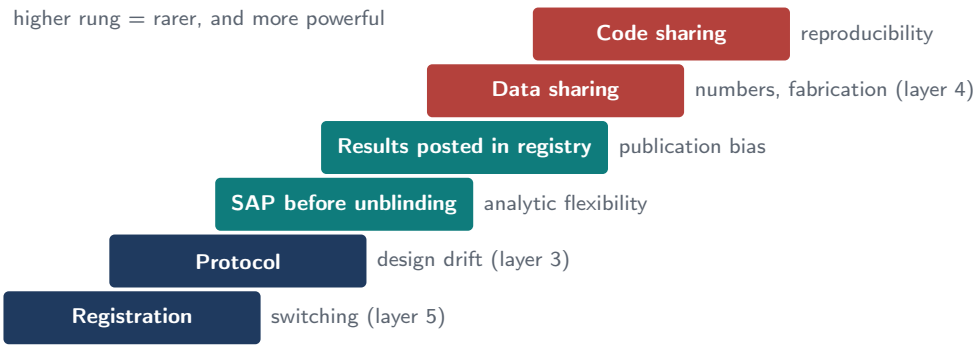


COMPare: all 67 trials over 6 weeks in NEJM, Lancet, JAMA, BMJ, Annals

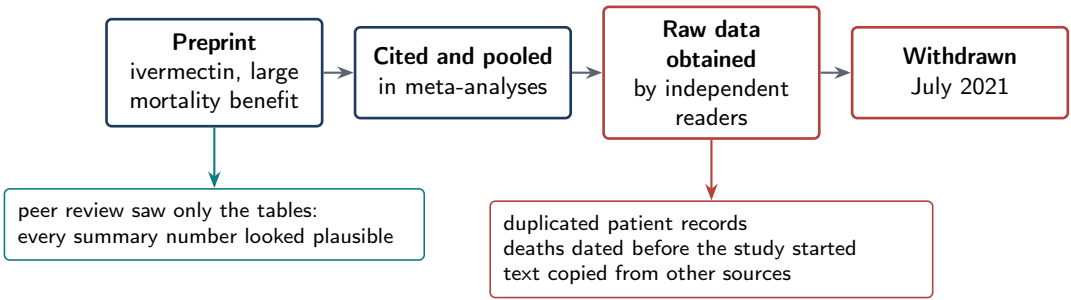


Goldacre et al. COMPare. Trials 2019;20:118 (materials/pdfs/method-goldacre-compare-2019.pdf).

Open science is a ladder: each rung makes one more check possible

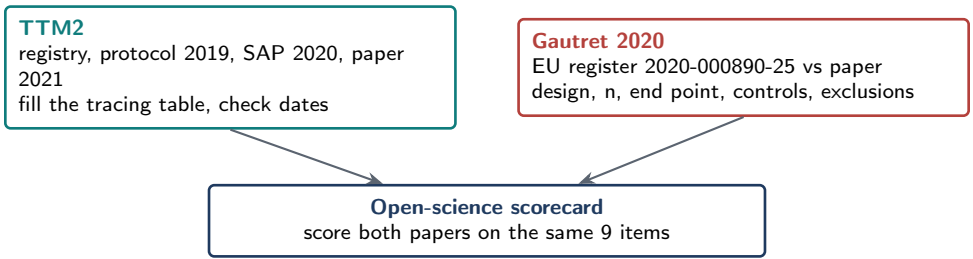


Elgazzar 2020: only the raw data could show it



Research Square withdrawal note, 14 July 2021 (materials/pdfs/article-elgazzar-2020-withdrawn.pdf).

Exercise 01 Trace two papers from promise to publication



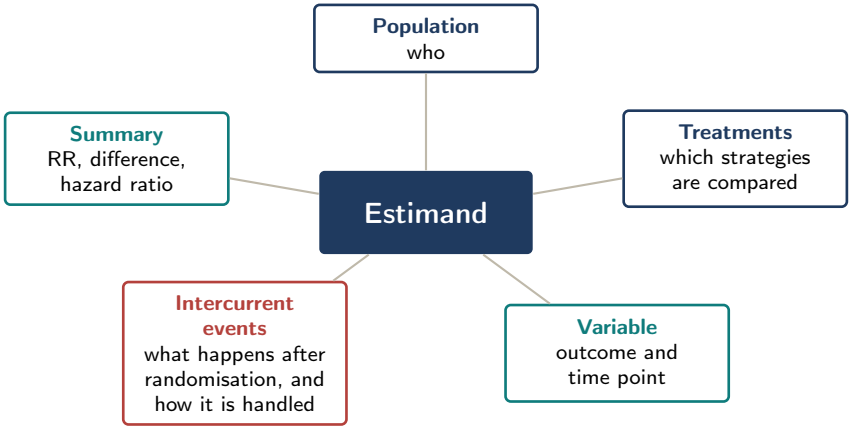
Pairs **60 minutes** **workshop/01-open-science.md**

Part A TTM2 (25 min), Part B Gautret (20 min), Part C COMPare (5 min), Part D scorecard (10 min).

2 Statistical methodology on real papers

Estimand, confidence intervals, forensics, bias, spin. Layers 2, 4, 6 and 7, on TTM2 and Gautret.

The estimand: five attributes that define the question



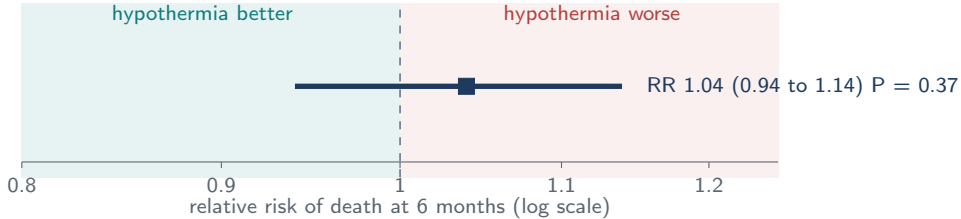
ICH E9(R1) addendum on estimands, 2019.

Same words, different questions: TTM2 vs Gautret

	TTM2	Gautret
Population	comatose adults after OHCA	inpatients, one centre; controls from other centres
Treatments	33 °C vs normothermia	HCQ (± AZ) vs no HCQ, not assigned
Variable	death by 6 months	PCR negative at day 6 (surrogate)
Intercurrent	treatment policy	6 treated excluded: ICU, death, discharge, stop
Summary	relative risk	proportions, many p-values

Dankiewicz et al. NEJM 2021; Gautret et al. Int J Antimicrob Agents 2020 (retracted 2024).

Read the interval, not the p-value

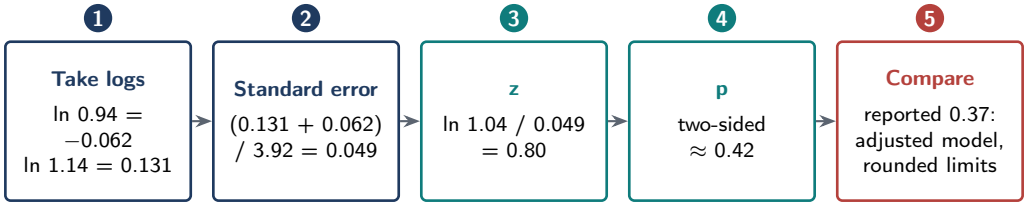


Correct
no difference shown; effects from a 6% reduction to a 14% increase remain compatible

Wrong
"hypothermia does not work"
"the two treatments are equivalent"

TTM2: death 465/925 (50%) vs 446/925 (48%).

From a confidence interval back to a p-value



Rule of thumb. A 95% CI that excludes the null implies $p < 0.05$; one that includes it implies $p > 0.05$. A paper that breaks this rule needs a question to the authors.

Altman and Bland, BMJ 2011;343:d2090. Toolkit, layer 4.

Forensics: four checks you can do with a phone

Percentages

recompute every n/N you can find

87/120 = 72.5%, not 75.5%

CI vs p

do they agree?

OR 1.32 (1.01 to 1.73) with p = 0.11: no

GRIM

mean × n must be close to an integer

n = 13, mean 2.41: 31.33: impossible

Flow

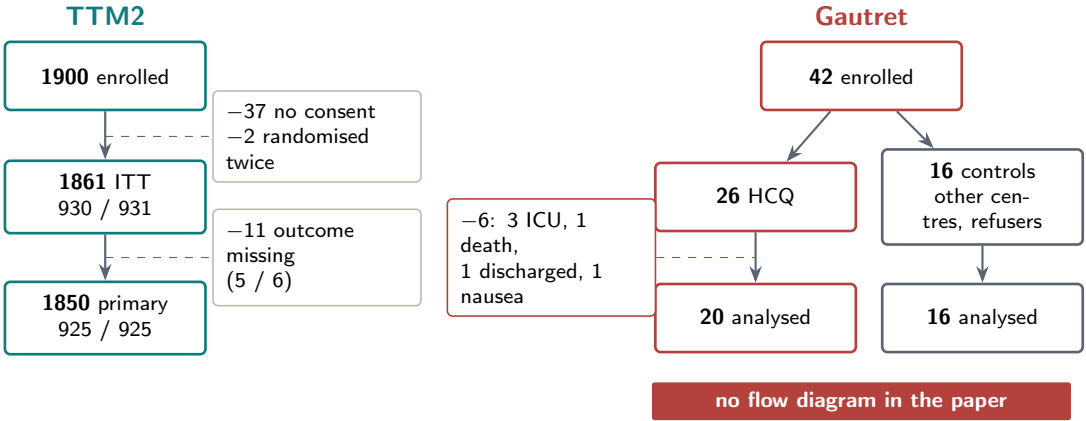
every patient from enrolled to analysed

248 randomised, 124 + 122 shown

an inconsistency is a question to the authors, not an accusation

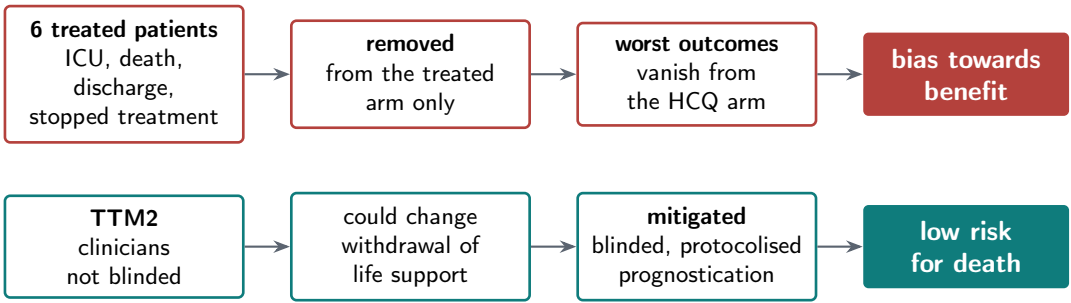
Brown and Heathers, Soc Psychol Personal Sci 2017 (GRIM); Nuijten et al., Behav Res Methods 2016 (statcheck).

Reconcile the flow: TTM2 adds up, Gautret does not show it



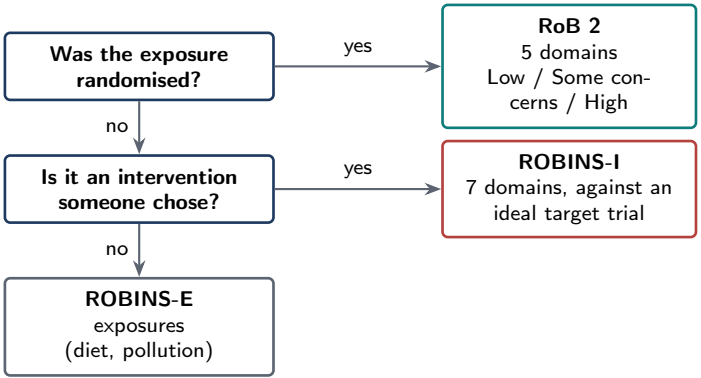
TTM2 abstract and Fig. S1; Gautret 2020, section 3.1.

Bias has a direction: ask which way it pushes



Calling outcome-dependent exclusions "per protocol" does not remove the bias.

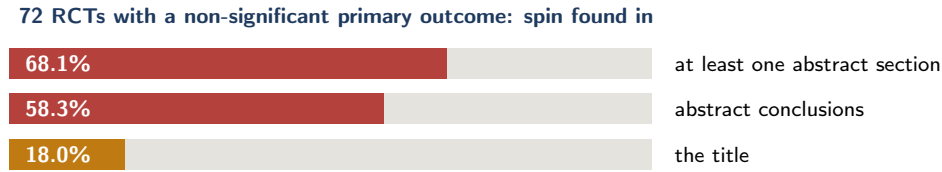
Which risk-of-bias tool?



one **Critical** domain makes the whole study Critical: Gautret (confounding, missing data)

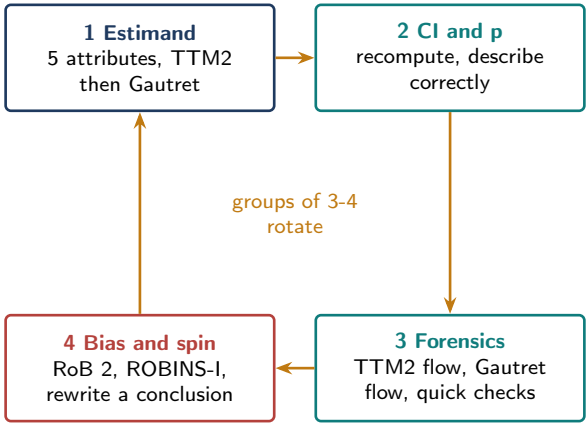
Sterne et al. BMJ 2019 (RoB 2); BMJ 2016 (ROBINS-I); Higgins et al. Environ Int 2024 (ROBINS-E).

Spin: the abstract says more than the results



Boutron et al. JAMA 2010;303:2058 (materials/pdfs-local/method-boutron-spin-2010.pdf).

Exercise 02 Four stations, 20 minutes each

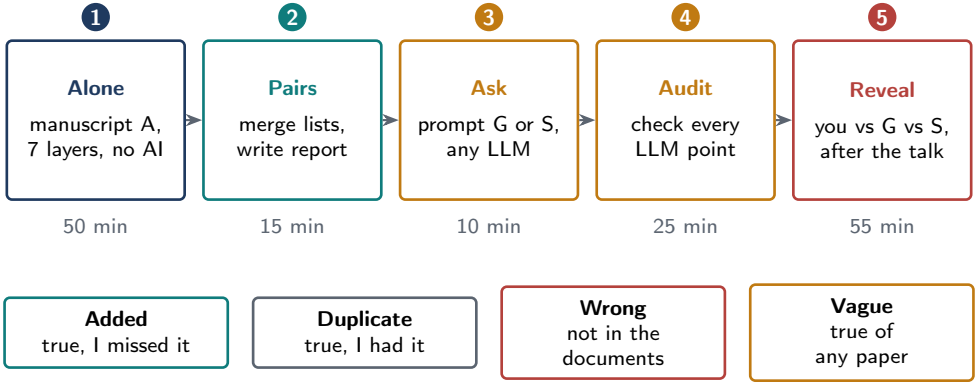


workshop/02-methodology-stations.md

3 Apply the method, then ask an LLM

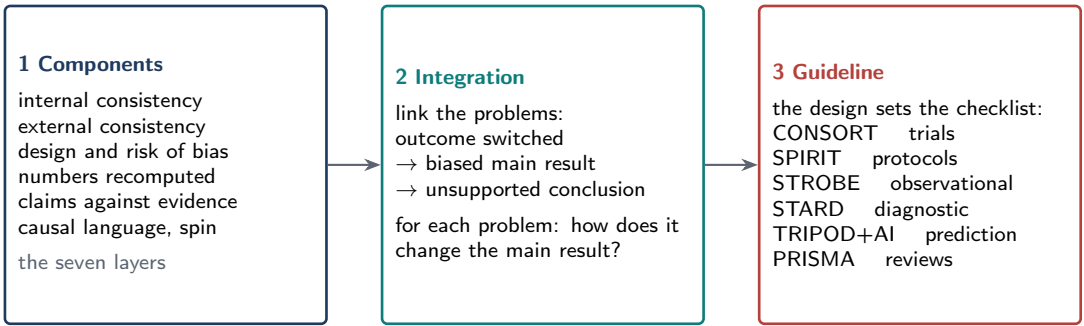
One new paper, reviewed alone with the seven layers. Then an LLM reviews it too, with a generic or a structured prompt.

How the afternoon works



workshop/03-review-and-llm.md; llm-exercise/manuscript-A.md. Manuscript A is fictional: never paste a manuscript you review for a journal.

A review is a process: components, integration, guideline



The same mental model is what you give an LLM in prompt S. Links: handouts/toolkit.md, layer 1.

Two prompts, same paper

Prompt G (generic), odd seats

You are a peer reviewer for a general medical journal. Review the paper below.

Prompt S (structured), even seats

Work through these components in order; each feeds the next:

1 article type and its reporting guideline
2 external consistency with every source given
3 internal consistency: recompute every number
4 design and risk of bias
5 claims and language
6 integration: how each problem changes the conclusion

Location, what is wrong (show arithmetic), why it matters. Rank. Invent nothing.

Both

Paste from "Document 1" to the end. Note the model and its version.

Full text of both prompts: workshop/03-review-and-llm.md, step 3.

The review report: what an author can act on

Summary 3 to 5 neutral sentences: you understood the paper

Major issues at most four, each with three parts → **1 the problem** **2 why it matters** **3 what to do**

Minor issues wording, tables, references

Recommendation and your confidence in the main conclusion

Confidential to the editor concerns you cannot prove

Not: "The statistics are poor." Instead: "Six treated patients are excluded after ICU transfer or death (Section 3.1). Because they are the worst outcomes, this biases the result towards benefit. Please report all 26 treated patients."

4 LLMs in peer review: practical concerns

Policies, confidentiality, invented facts, and what the evidence shows.
Then the reveal: you, the generic prompt and the structured prompt.

Where an LLM could enter, and where it must not

Your own paper
before submission:
check reporting

Published papers
screening, extraction
(verified)

Confidential manuscript
you review for
a journal

Your written report
language polishing

allowed, with verification and disclosure

do not upload

only if policy allows

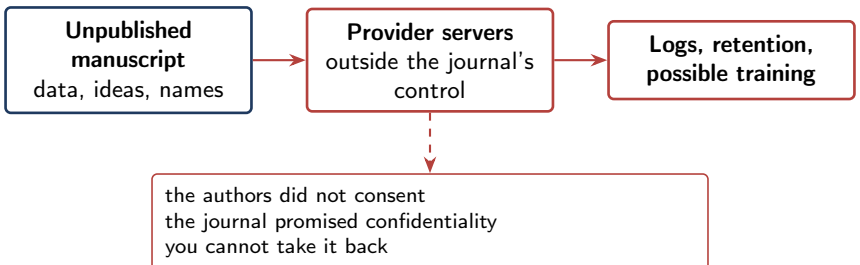
NIH (2023)
reviewers may not use
generative AI on applications

ICMJE
no upload where confidentiality
cannot be assured; disclose use

Most publishers
reviewers must not upload
manuscripts to AI tools

NIH notice NOT-OD-23-149 (2023); ICMJE Recommendations (updated 2024); check your journal's current policy before every review.

Confidentiality: what happens when you paste a manuscript



Today's exercise is safe because manuscript A is fictional and written for training.

Four ways an LLM review goes wrong

Invented facts

registry content, protocol details it never saw

Invented refs

plausible authors, journals, years; no such paper

Wrong arithmetic

percentages, p from CI, GRIM, stated confidently

Boilerplate

"consider confounding" for any paper

every LLM claim is a hypothesis you must check in the document

What the evidence shows so far

GPT-4 feedback on 3096 Nature-family papers (Liang et al.)



LLM RoB 2 judgements vs Cochrane reviews, 40 RCTs (Huang et al.)



Liang et al. NEJM AI 2024; Huang et al. J Med Internet Res 2025;27:e70450. Useful as a second reader; not a replacement.

When AI conducts research: reproducible is not valid

PNAS 2026: 288 researchers randomised in 103 teams to reproduce published results



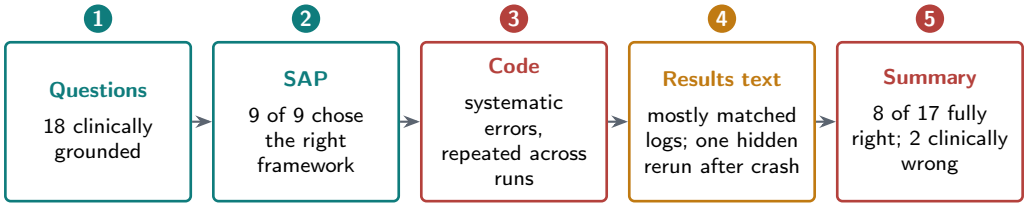
REPRO-Bench
ACL 2025
best agent: **21.4%** on
112 reproducibility tasks

AI scientist pitfalls
arXiv 2025
benchmark choice, leakage,
metric misuse,
post-hoc selection

AI Scientist
Nature 2026
its paper passed first-round
review at an ML workshop

Brodeur et al. PNAS 2026 (social-science reproduction only); Hu et al. Findings ACL 2025; Luo, Kasirzadeh, Shah, arXiv:2509.08713; Lu et al. Nature 2026.

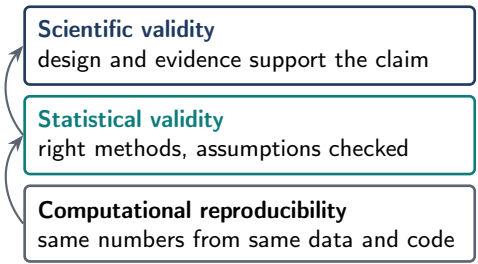
An LLM agent on a clinical survival analysis



Lesson. A good plan did not predict correct code. Experts still had to check formulas, cohort boundaries and the concordance computation before any result could be reported.

Wu et al. J Med Internet Res 2026 (doi:10.2196/99597): Claude agent, public Moorfields AMD dataset (7802 eyes), Kaplan-Meier and Cox, 27 runs.

The code runs. Is the conclusion valid?

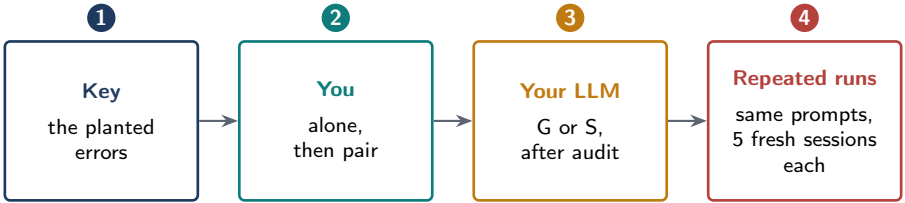


Case. An agent compares survival on treatment A vs B: Cox HR 0.72, $p < 0.05$, no errors.

Check: randomised or confounded? time zero (immortal time)? proportional hazards? missing data? does a hazard ratio answer the clinical question? SAP fixed before results?

Each level needs the one below it, but does not follow from it.

Reveal and compare



	manuscript only	+ registry and SAP
prompt G	?	? (half the room)
prompt S	?	? (half the room)

Errors found, wrong claims, traps flagged. What changes with the prompt? What changes with the sources?

What repeated runs teach, for any paper

Documents beat prompts

a check that needs the registry or SAP cannot be done without it, whatever the prompt

Structure raises the right questions

the structured prompt asked what it could not verify; the generic one mostly did not

Arithmetic is cheap, not free

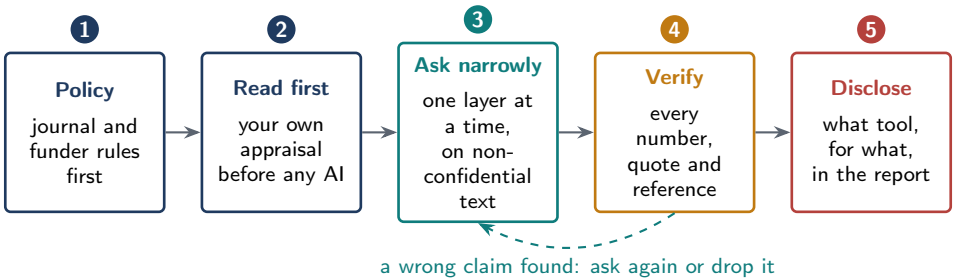
recomputation was near ceiling, yet some runs saw an odd number and did not test it

More text, more to audit

structured reviews were 60 to 80% longer; every point still needs checking

The structure came from the method you learned this morning and from the reporting guideline. Clean planted errors are an easy case: real manuscripts are harder, and a real manuscript under review never goes into an LLM.

A safe routine for AI in your own appraisal work



Wrap-up

Open science
compare the paper with what was promised

Methodology
estimand, interval, numbers, bias direction, spin

Review
problem, why it matters, what to do

LLMs
never confidential text; verify every claim; disclose

Post-quiz, then your morning sheet

What would you now check in the TTM2 abstract that you did not check this morning? Has your confidence in its conclusion changed?

Take-home kit: 7-layer checklist, toolkit, review report template. To go further: the reporting guideline for the designs you read (toolkit, layer 1).